

# Srgan

November 21, 2023

## 0.0.1 Libraries

```
[1]: import numpy as np
import pandas as pd
import os, math, sys
import glob, itertools
import argparse, random

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable
from torchvision.models import vgg19
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, Dataset
from torchvision.utils import save_image, make_grid

import plotly
import plotly.express as px
import plotly.graph_objects as go
import matplotlib.pyplot as plt

from PIL import Image
from tqdm import tqdm_notebook as tqdm
from sklearn.model_selection import train_test_split

random.seed(42)
import warnings
warnings.filterwarnings("ignore")
```

## 0.0.2 Settings

```
[2]: load_pretrained_models = True
n_epochs = 2
dataset_path = "../input/celeba-dataset/img_align_celeba/img_align_celeba"
batch_size = 16
lr = 0.00008
b1 = 0.5
```

```

b2 = 0.999
decay_epoch = 100
n_cpu = 8
hr_height = 256
hr_width = 256
channels = 3

os.makedirs("images", exist_ok=True)
os.makedirs("saved_models", exist_ok=True)

cuda = torch.cuda.is_available()
hr_shape = (hr_height, hr_width)

```

### 0.0.3 Define Dataset Class

```

[3]: mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])

class ImageDataset(Dataset):
    def __init__(self, files, hr_shape):
        hr_height, hr_width = hr_shape
        self.lr_transform = transforms.Compose(
            [
                transforms.Resize((hr_height // 4, hr_height // 4), Image.
                    BICUBIC),
                transforms.ToTensor(),
                transforms.Normalize(mean, std),
            ]
        )
        self.hr_transform = transforms.Compose(
            [
                transforms.Resize((hr_height, hr_height), Image.BICUBIC),
                transforms.ToTensor(),
                transforms.Normalize(mean, std),
            ]
        )
        self.files = files

    def __getitem__(self, index):
        img = Image.open(self.files[index % len(self.files)])
        img_lr = self.lr_transform(img)
        img_hr = self.hr_transform(img)

        return {"lr": img_lr, "hr": img_hr}

    def __len__(self):
        return len(self.files)

```

## 0.0.4 Get Train/Test Dataloaders

```
[5]: train_paths, test_paths = train_test_split(sorted(glob.glob(dataset_path + "/*.*")), test_size=0.02, random_state=42)
train_dataloader = DataLoader(ImageDataset(train_paths, hr_shape=hr_shape), batch_size=batch_size, shuffle=True, num_workers=n_cpu)
test_dataloader = DataLoader(ImageDataset(test_paths, hr_shape=hr_shape), batch_size=int(batch_size*0.75), shuffle=True, num_workers=n_cpu)
```

## 0.0.5 Define Model Classes

```
[15]: class FeatureExtractor(nn.Module):
    def __init__(self):
        super(FeatureExtractor, self).__init__()
        vgg19_model = vgg19(pretrained=True)
        self.feature_extractor = nn.Sequential(*list(vgg19_model.features.children())[:18])

    def forward(self, img):
        return self.feature_extractor(img)

class ResidualBlock(nn.Module):
    def __init__(self, in_features):
        super(ResidualBlock, self).__init__()
        self.conv_block = nn.Sequential(
            nn.Conv2d(in_features, in_features, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(in_features, 0.8),
            nn.PReLU(),
            nn.Conv2d(in_features, in_features, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(in_features, 0.8),
        )

    def forward(self, x):
        return x + self.conv_block(x)

class GeneratorResNet(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, n_residual_blocks=16):
        super(GeneratorResNet, self).__init__()

        # First layer
        self.conv1 = nn.Sequential(nn.Conv2d(in_channels, 64, kernel_size=9, stride=1, padding=4), nn.PReLU())
```

```

# Residual blocks
res_blocks = []
for _ in range(n_residual_blocks):
    res_blocks.append(ResidualBlock(64))
self.res_blocks = nn.Sequential(*res_blocks)

# Second conv layer post residual blocks
self.conv2 = nn.Sequential(nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1), nn.BatchNorm2d(64, 0.8))

# Upsampling layers
upsampling = []
for out_features in range(2):
    upsampling += [
        # nn.Upsample(scale_factor=2),
        nn.Conv2d(64, 256, 3, 1, 1),
        nn.BatchNorm2d(256),
        nn.PixelShuffle(upscale_factor=2),
        nn.PReLU(),
    ]
self.upsampling = nn.Sequential(*upsampling)

# Final output layer
self.conv3 = nn.Sequential(nn.Conv2d(64, out_channels, kernel_size=9, stride=1, padding=4), nn.Tanh())

def forward(self, x):
    out1 = self.conv1(x)
    out = self.res_blocks(out1)
    out2 = self.conv2(out)
    out = torch.add(out1, out2)
    out = self.upsampling(out)
    out = self.conv3(out)
    return out

class Discriminator(nn.Module):
    def __init__(self, input_shape):
        super(Discriminator, self).__init__()

        self.input_shape = input_shape
        in_channels, in_height, in_width = self.input_shape
        patch_h, patch_w = int(in_height / 2 ** 4), int(in_width / 2 ** 4)
        self.output_shape = (1, patch_h, patch_w)

    def discriminator_block(in_filters, out_filters, first_block=False):
        layers = []

```

```

        layers.append(nn.Conv2d(in_filters, out_filters, kernel_size=3, u
˓→stride=1, padding=1))
        if not first_block:
            layers.append(nn.BatchNorm2d(out_filters))
        layers.append(nn.LeakyReLU(0.2, inplace=True))
        layers.append(nn.Conv2d(out_filters, out_filters, kernel_size=3, u
˓→stride=2, padding=1))
        layers.append(nn.BatchNorm2d(out_filters))
        layers.append(nn.LeakyReLU(0.2, inplace=True))
    return layers

layers = []
in_filters = in_channels
for i, out_filters in enumerate([64, 128, 256, 512]):
    layers.extend(discriminator_block(in_filters, out_filters, u
˓→first_block=(i == 0)))
    in_filters = out_filters

layers.append(nn.Conv2d(out_filters, 1, kernel_size=3, stride=1, u
˓→padding=1))

self.model = nn.Sequential(*layers)

def forward(self, img):
    return self.model(img)

```

## 0.0.6 Train Super Resolution GAN (SRGAN)

```
[17]: # Initializing generator and discriminator
generator = GeneratorResNet()
discriminator = Discriminator(input_shape=(channels, *hr_shape))
feature_extractor = FeatureExtractor()

# Setting feature extractor to inference mode
feature_extractor.eval()

# Losses calculation
criterion_GAN = torch.nn.MSELoss()
criterion_content = torch.nn.L1Loss()

if cuda:
    generator = generator.cuda()
    discriminator = discriminator.cuda()
    feature_extractor = feature_extractor.cuda()
    criterion_GAN = criterion_GAN.cuda()
    criterion_content = criterion_content.cuda()
```

```

# Load pretrained models
if load_pretrained_models:
    generator.load_state_dict(torch.load("../input/
↪single-image-super-resolution-gan-srgan-pytorch/saved_models/generator.pth"))
    discriminator.load_state_dict(torch.load("../input/
↪single-image-super-resolution-gan-srgan-pytorch/saved_models/discriminator.
↪pth"))

# Optimizers
optimizer_G = torch.optim.Adam(generator.parameters(), lr=lr, betas=(b1, b2))
optimizer_D = torch.optim.Adam(discriminator.parameters(), lr=lr, betas=(b1, ↪
↪b2))

Tensor = torch.cuda.FloatTensor if cuda else torch.Tensor

```

```

[18]: train_gen_losses, train_disc_losses, train_counter = [], [], []
test_gen_losses, test_disc_losses = [], []
test_counter = [idx*len(train_dataloader.dataset) for idx in range(1, ↪
↪n_epochs+1)]

for epoch in range(n_epochs):

    ### Training
    gen_loss, disc_loss = 0, 0
    tqdm_bar = tqdm(train_dataloader, desc=f'Training Epoch {epoch} ', ↪
↪total=int(len(train_dataloader)))
    for batch_idx, imgs in enumerate(tqdm_bar):
        generator.train(); discriminator.train()
        # Configure model input
        imgs_lr = Variable(imgs["lr"].type(Tensor))
        imgs_hr = Variable(imgs["hr"].type(Tensor))
        # Adversarial ground truths
        valid = Variable(Tensor(np.ones((imgs_lr.size(0), *discriminator.
↪output_shape))), requires_grad=False)
        fake = Variable(Tensor(np.zeros((imgs_lr.size(0), *discriminator.
↪output_shape))), requires_grad=False)

        ### Train Generator
        optimizer_G.zero_grad()
        # Generate a high resolution image from low resolution input
        gen_hr = generator(imgs_lr)
        # Adversarial loss
        loss_GAN = criterion_GAN(discriminator(gen_hr), valid)
        # Content loss
        gen_features = feature_extractor(gen_hr)

```

```

real_features = feature_extractor(imgs_hr)
loss_content = criterion_content(gen_features, real_features.detach())
# Total loss
loss_G = loss_content + 1e-3 * loss_GAN
loss_G.backward()
optimizer_G.step()

#### Train Discriminator
optimizer_D.zero_grad()
# Loss of real and fake images
loss_real = criterion_GAN(discriminator(imgs_hr), valid)
loss_fake = criterion_GAN(discriminator(gen_hr.detach()), fake)
# Total loss
loss_D = (loss_real + loss_fake) / 2
loss_D.backward()
optimizer_D.step()

gen_loss += loss_G.item()
train_gen_losses.append(loss_G.item())
disc_loss += loss_D.item()
train_disc_losses.append(loss_D.item())
train_counter.append(batch_idx*batch_size + imgs_lr.size(0) + epoch*len(train_dataloader.dataset))
tqdm_bar.set_postfix(gen_loss=gen_loss/(batch_idx+1), disc_loss=disc_loss/(batch_idx+1))

# Testing
gen_loss, disc_loss = 0, 0
tqdm_bar = tqdm(test_dataloader, desc=f'Testing Epoch {epoch} ', total=int(len(test_dataloader)))
for batch_idx, imgs in enumerate(tqdm_bar):
    generator.eval(); discriminator.eval()
    # Configure model input
    imgs_lr = Variable(imgs["lr"].type(Tensor))
    imgs_hr = Variable(imgs["hr"].type(Tensor))
    # Adversarial ground truths
    valid = Variable(Tensor(np.ones((imgs_lr.size(0), *discriminator.output_shape))), requires_grad=False)
    fake = Variable(Tensor(np.zeros((imgs_lr.size(0), *discriminator.output_shape))), requires_grad=False)

    #### Eval Generator
    # Generate a high resolution image from low resolution input
    gen_hr = generator(imgs_lr)
    # Adversarial loss
    loss_GAN = criterion_GAN(discriminator(gen_hr), valid)
    # Content loss

```

```

gen_features = feature_extractor(gen_hr)
real_features = feature_extractor(imgs_hr)
loss_content = criterion_content(gen_features, real_features.detach())
# Total loss
loss_G = loss_content + 1e-3 * loss_GAN

### Eval Discriminator
# Loss of real and fake images
loss_real = criterion_GAN(discriminator(imgs_hr), valid)
loss_fake = criterion_GAN(discriminator(gen_hr.detach()), fake)
# Total loss
loss_D = (loss_real + loss_fake) / 2

gen_loss += loss_G.item()
disc_loss += loss_D.item()
tqdm_bar.set_postfix(gen_loss=gen_loss/(batch_idx+1), disc_loss=disc_loss/(batch_idx+1))

# Saving image grid with upsampled inputs and SRGAN outputs
if random.uniform(0,1)<0.1:
    imgs_lr = nn.functional.interpolate(imgs_lr, scale_factor=4)
    imgs_hr = make_grid(imgs_hr, nrow=1, normalize=True)
    gen_hr = make_grid(gen_hr, nrow=1, normalize=True)
    imgs_lr = make_grid(imgs_lr, nrow=1, normalize=True)
    img_grid = torch.cat((imgs_hr, imgs_lr, gen_hr), -1)
    save_image(img_grid, f"images/{batch_idx}.png", normalize=False)

test_gen_losses.append(gen_loss/len(test_dataloader))
test_disc_losses.append(disc_loss/len(test_dataloader))

# Saving model checkpoints
if np.argmax(test_gen_losses) == len(test_gen_losses)-1:
    torch.save(generator.state_dict(), "saved_models/generator.pth")
    torch.save(discriminator.state_dict(), "saved_models/discriminator.pth")

HBox(children=(FloatProgress(value=0.0, description='Training Epoch 0 ', max=12410.0, style=ProgressStyle(desc...))

HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 0 ', max=338.0, style=ProgressStyle(descr...))

HBox(children=(FloatProgress(value=0.0, description='Training Epoch 1 ', max=12410.0, style=ProgressStyle(desc...

```

```
HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 1 ', max=338.  
↳0, style=ProgressStyle(descrip...
```

```
[19]: import plotly.graph_objects as go  
  
fig = go.Figure()  
  
# Train Generator Loss  
fig.add_trace(go.Scatter(x=train_counter, y=train_gen_losses, mode='lines',  
↳name='Train Generator Loss'))  
  
# Test Generator Loss  
fig.add_trace(go.Scatter(x=test_counter, y=test_gen_losses,  
                        mode='markers', marker=dict(symbol='star-diamond',  
↳color='orange', size=9, line_width=1),  
                        name='Test Generator Loss'))  
  
fig.update_layout(  
    width=1000,  
    height=500,  
    title="Train vs. Test Generator Loss",  
    xaxis_title="Number of training examples seen",  
    yaxis_title="Adversarial + Content Loss",  
    legend=dict(x=0, y=1, traceorder='normal', orientation='h'),  
)  
  
fig.show()
```

```
[20]: import plotly.graph_objects as go  
  
fig = go.Figure()  
  
# Train Discriminator Loss  
fig.add_trace(go.Scatter(x=train_counter, y=train_disc_losses, mode='lines',  
↳name='Train Discriminator Loss'))  
  
# Test Discriminator Loss  
fig.add_trace(go.Scatter(x=test_counter, y=test_disc_losses,  
                        mode='markers', marker=dict(symbol='star-diamond',  
↳color='orange', size=9, line_width=1),  
                        name='Test Discriminator Loss'))  
  
fig.update_layout(  
    width=1000,  
    height=500,
```

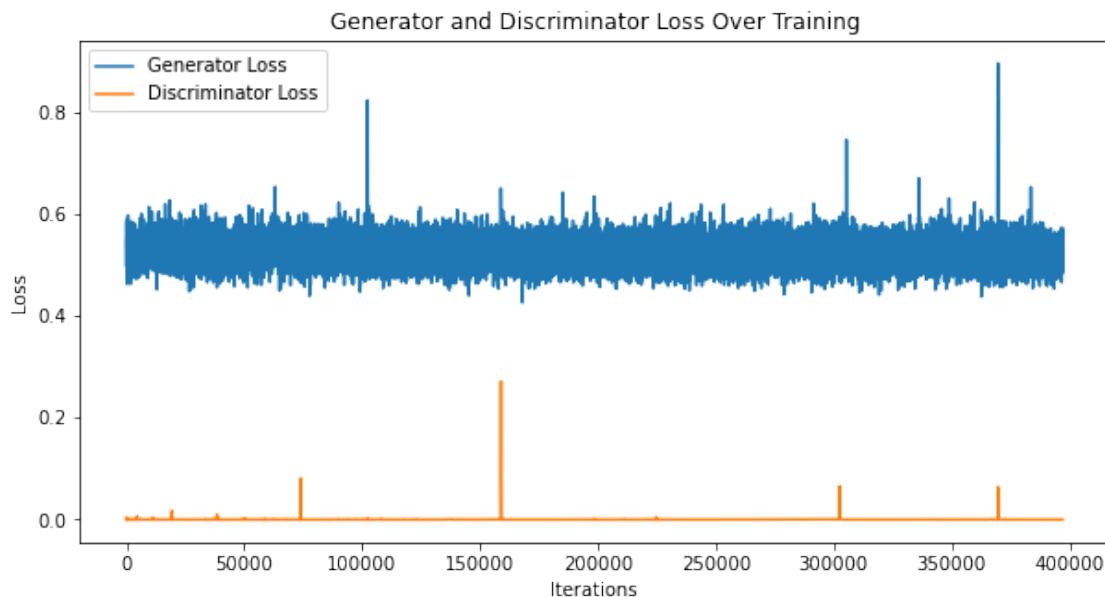
```

        title="Train vs. Test Discriminator Loss",
        xaxis_title="Number of training examples seen",
        yaxis_title="Adversarial Loss",
        legend=dict(x=0, y=1, traceorder='normal', orientation='h'),
    )

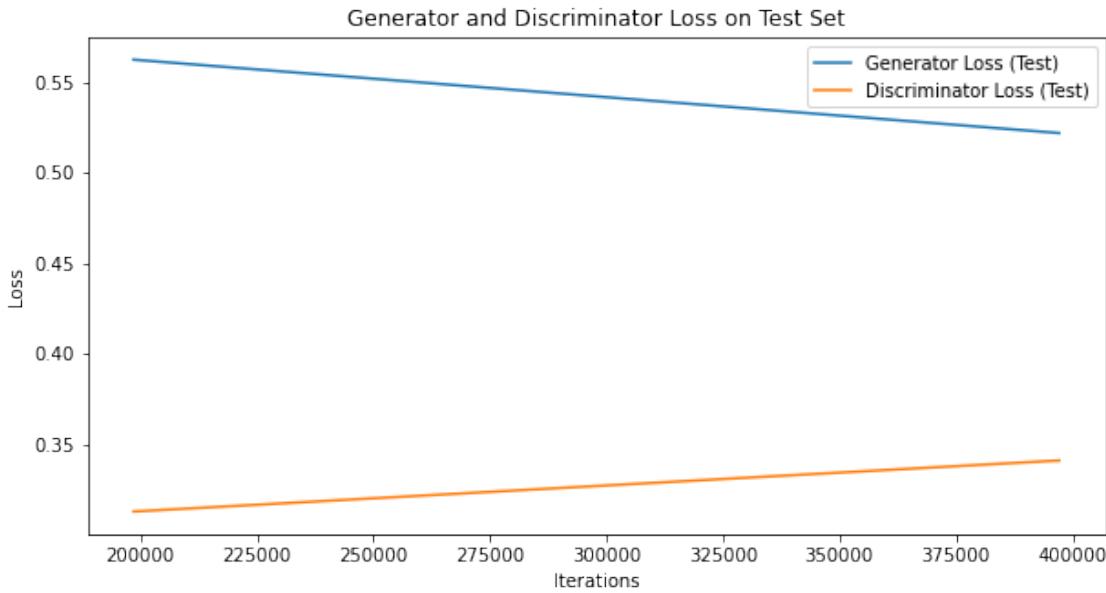
fig.show()

```

```
[28]: plt.figure(figsize=(10, 5))
plt.plot(train_counter, train_gen_losses, label='Generator Loss')
plt.plot(train_counter, train_disc_losses, label='Discriminator Loss')
plt.legend()
plt.title('Generator and Discriminator Loss Over Training')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.show()
```



```
[29]: plt.figure(figsize=(10, 5))
plt.plot(test_counter, test_gen_losses, label='Generator Loss (Test)')
plt.plot(test_counter, test_disc_losses, label='Discriminator Loss (Test)')
plt.legend()
plt.title('Generator and Discriminator Loss on Test Set')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.show()
```



```
[53]: from skimage.metrics import peak_signal_noise_ratio
import pandas as pd

psnr_details = []

for epoch in range(n_epochs):
    tqdm_bar = tqdm(test_dataloader, desc=f'Testing Epoch {epoch} ',  

                     total=int(len(test_dataloader)))
    for batch_idx, imgs in enumerate(tqdm_bar):
        generator.eval(); discriminator.eval()
        # Configure model input
        imgs_lr = Variable(imgs["lr"].type(Tensor))
        imgs_hr = Variable(imgs["hr"].type(Tensor))
        # Adversarial ground truths
        valid = Variable(Tensor(np.ones((imgs_lr.size(0), *discriminator.  

                                         output_shape))), requires_grad=False)
        fake = Variable(Tensor(np.zeros((imgs_lr.size(0), *discriminator.  

                                         output_shape))), requires_grad=False)

        ### Eval Generator
        # Generate a high resolution image from low resolution input
        gen_hr = generator(imgs_lr)
        # Adversarial loss
        loss_GAN = criterion_GAN(discriminator(gen_hr), valid)
        # Content loss
        gen_features = feature_extractor(gen_hr)
        real_features = feature_extractor(imgs_hr)

        psnr_details.append(peak_signal_noise_ratio(gen_hr, imgs_hr))
```

```

loss_content = criterion_content(gen_features, real_features.detach())
# Total loss
loss_G = loss_content + 1e-3 * loss_GAN

### Eval Discriminator
# Loss of real and fake images
loss_real = criterion_GAN(discriminator(imgs_hr), valid)
loss_fake = criterion_GAN(discriminator(gen_hr.detach()), fake)
# Total loss
loss_D = (loss_real + loss_fake) / 2

tqdm_bar.set_postfix(gen_loss=loss_G.item(), disc_loss=loss_D.item())

# Ensure images are in the range [0, 1] before calculating PSNR
imgs_hr = imgs_hr.clamp(0, 1)
gen_hr = gen_hr.clamp(0, 1)

# Calculate PSNR with data_range=1.0
psnr = peak_signal_noise_ratio(imgs_hr.cpu().detach().numpy(), gen_hr.
cpu().detach().numpy(), data_range=1.0)

# Store details in the list
psnr_details.append({
    'Epoch': epoch,
    'Batch': batch_idx,
    'PSNR': psnr,
    'HR_Min': imgs_hr.min().item(),
    'HR_Max': imgs_hr.max().item(),
    'Gen_Min': gen_hr.min().item(),
    'Gen_Max': gen_hr.max().item()
})

if random.uniform(0, 1) < 0.1:
    # Choose a random image index
    random_index = random.randint(0, len(test_paths) - 1)

# Bicubic upsampling
    imgs_lr_bicubic = nn.functional.interpolate(imgs_lr, scale_factor=4, mode='bicubic')

# Display the low-resolution, high-resolution, bicubic-upsampled, and generated images
plt.figure(figsize=(20, 5))

    plt.subplot(1, 4, 1)
    plt.imshow(transforms.ToPILImage()(imgs_lr.cpu()[0]))
    plt.title('Low Resolution')

```

```

plt.subplot(1, 4, 2)
plt.imshow(transforms.ToPILImage()(imgs_hr.cpu()[0]))
plt.title('High Resolution')

plt.subplot(1, 4, 3)
plt.imshow(transforms.ToPILImage()(imgs_lr_bicubic.cpu()[0]))
plt.title('Bicubic Upsampled')

plt.subplot(1, 4, 4)
plt.imshow(transforms.ToPILImage()(gen_hr.cpu()[0]))
plt.title('Generated (SRGAN)')

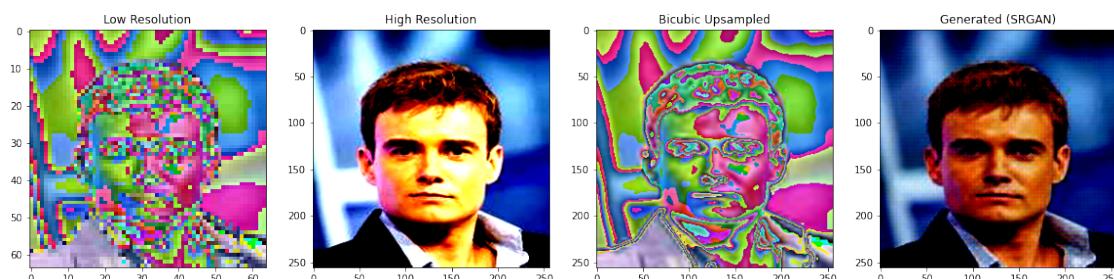
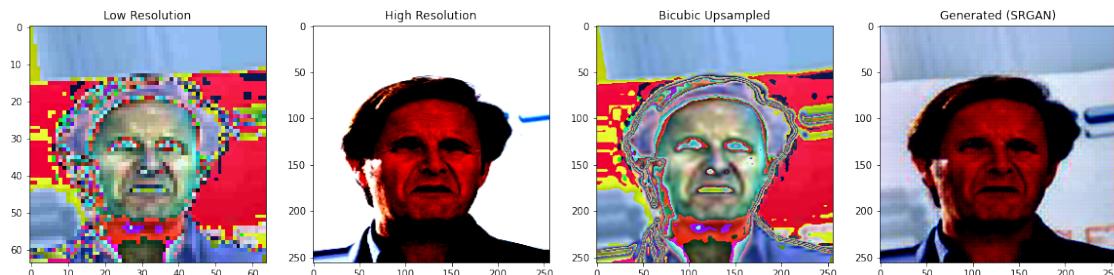
plt.show()

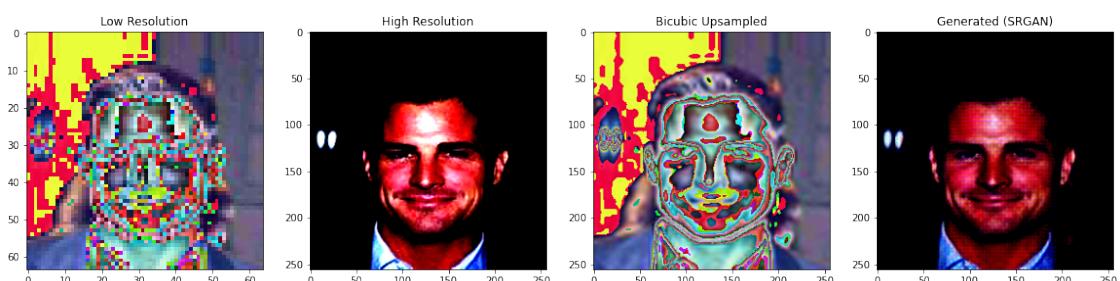
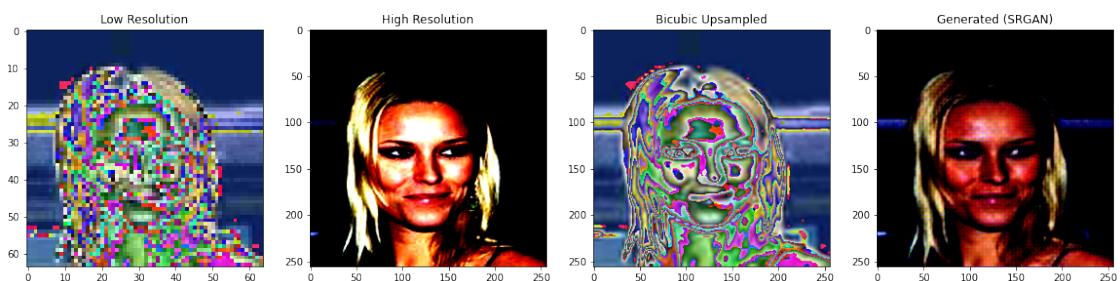
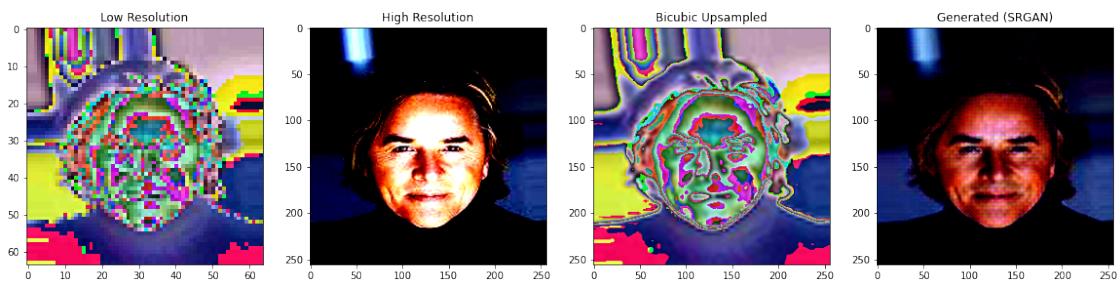
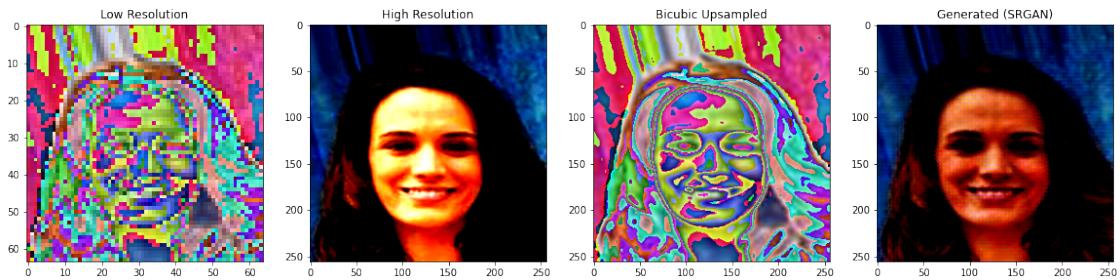
# Convert the list of dictionaries to a Pandas DataFrame
psnr_df = pd.DataFrame(psnr_details)

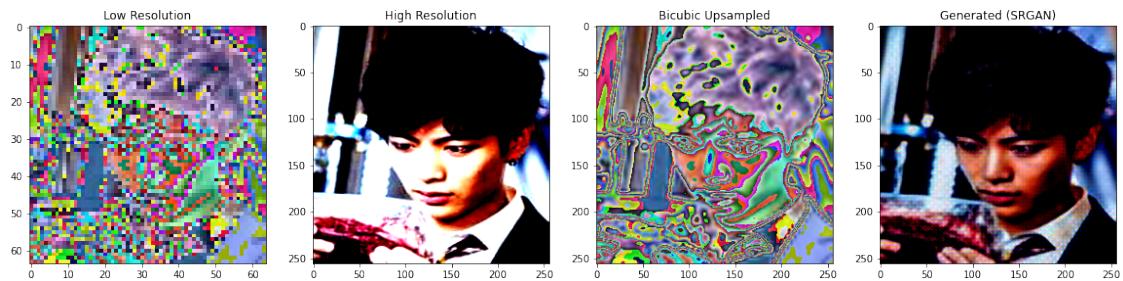
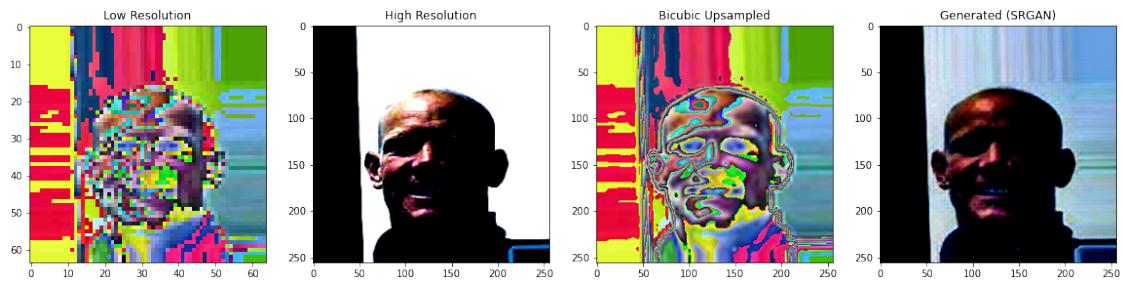
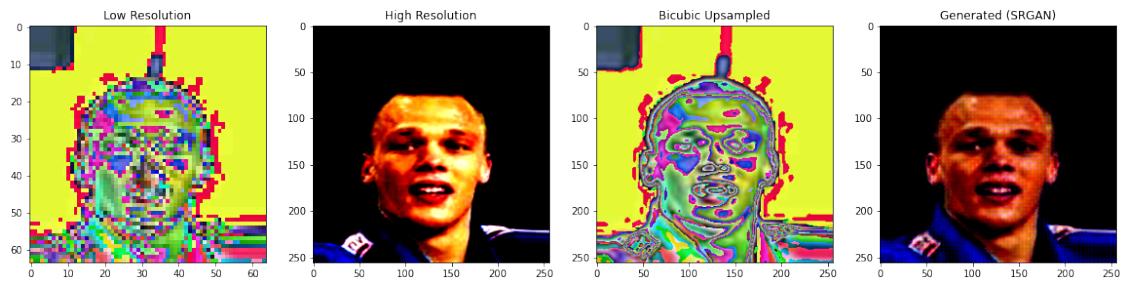
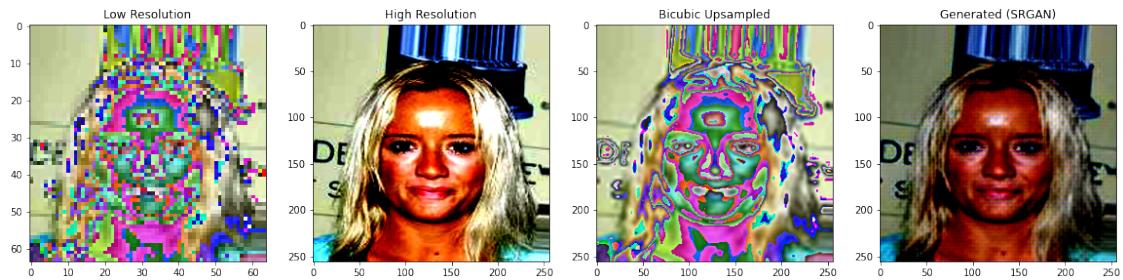
# Display the Pandas DataFrame
print(psnr_df)

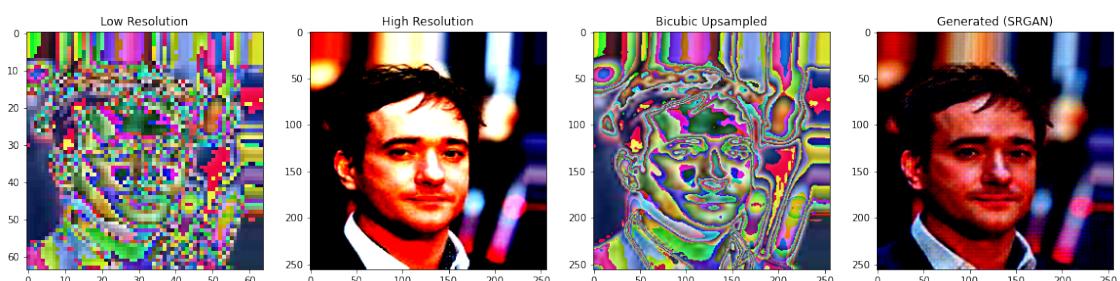
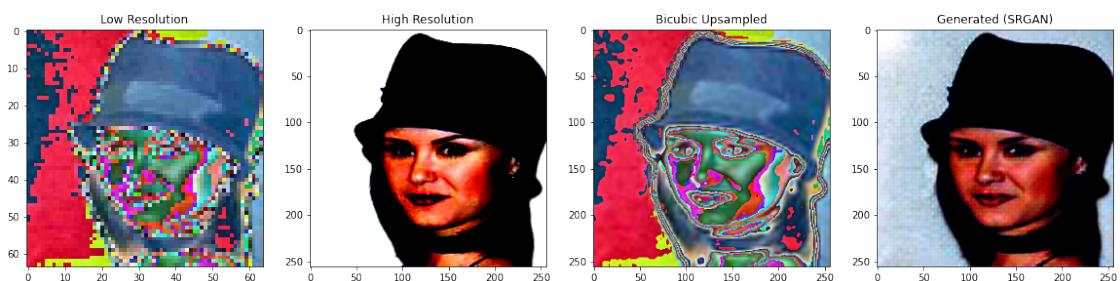
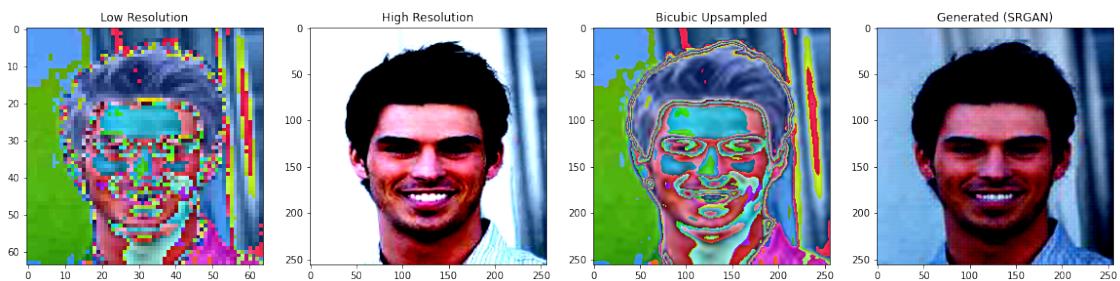
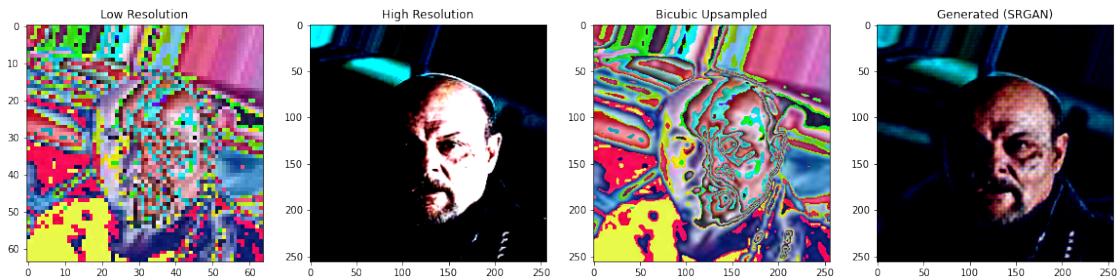
```

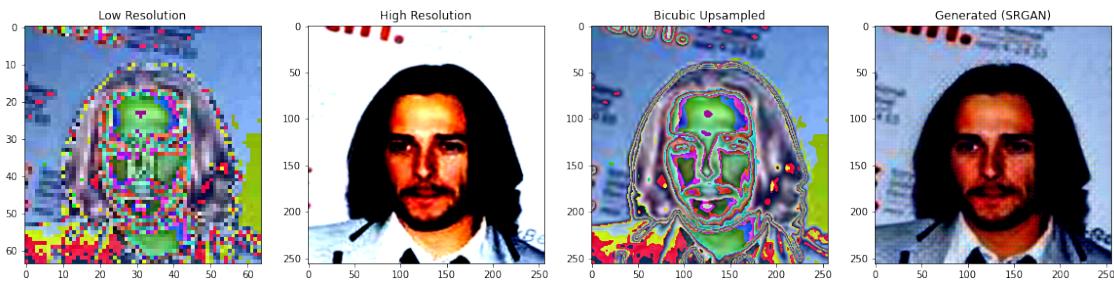
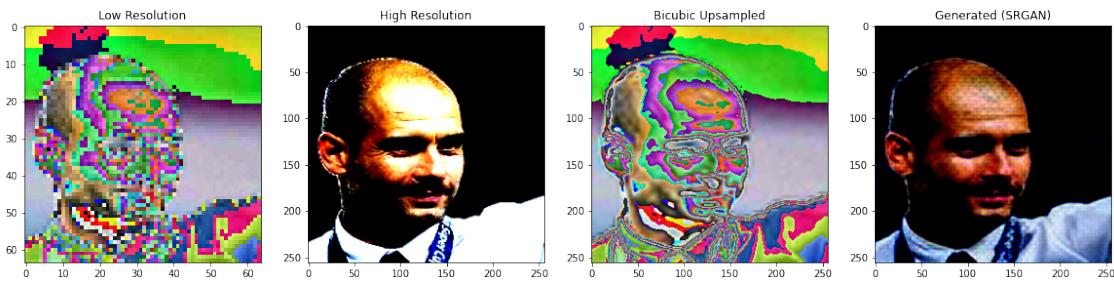
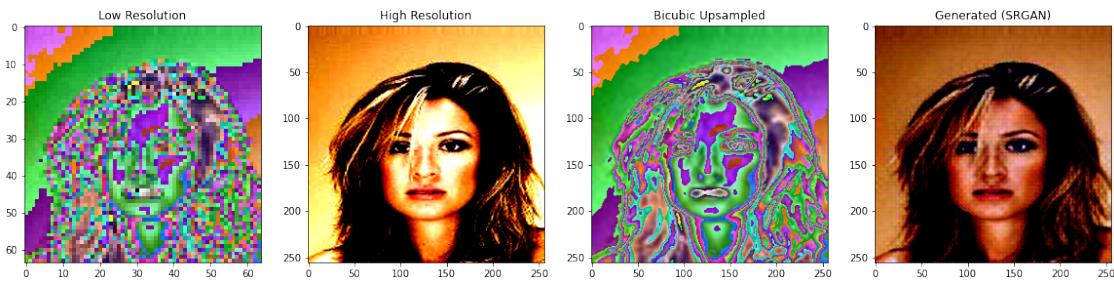
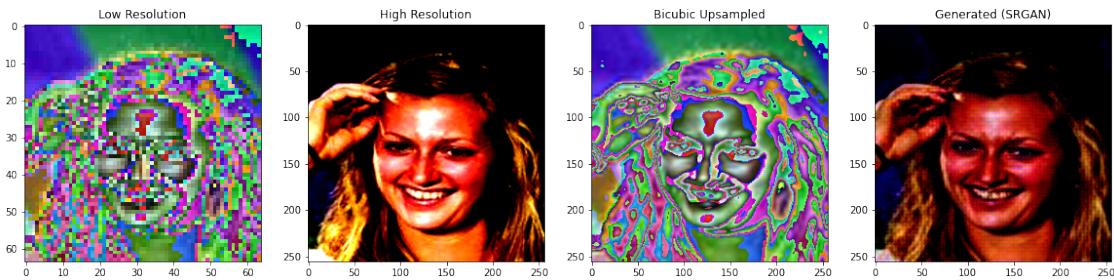
HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 0 ', max=338.  
`0, style=ProgressStyle(descrip...

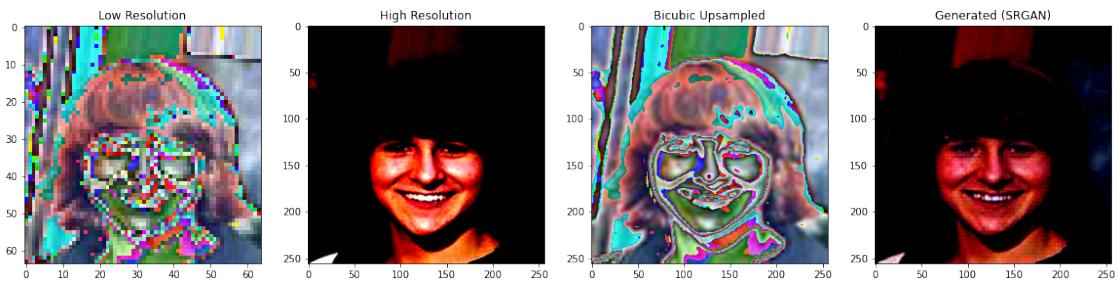
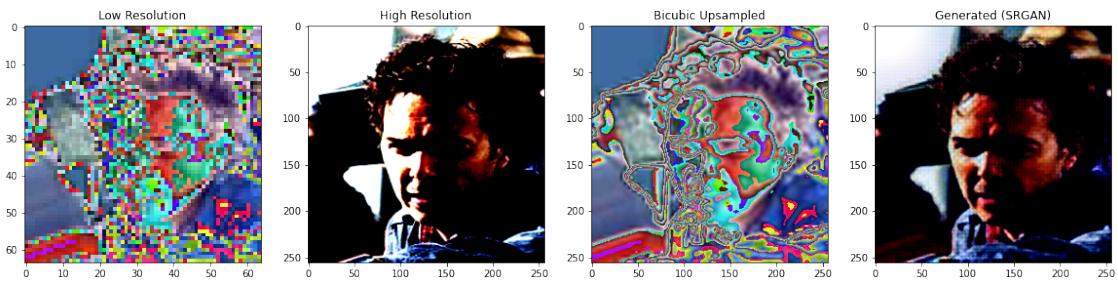
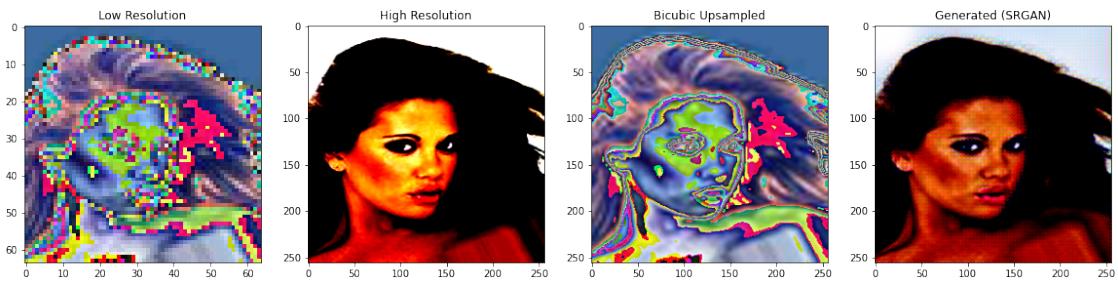
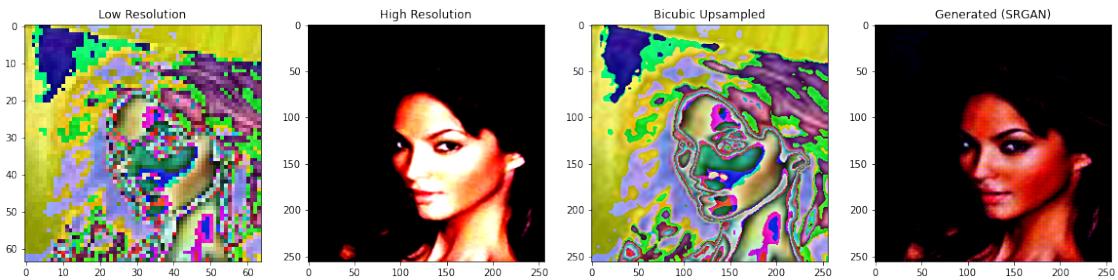


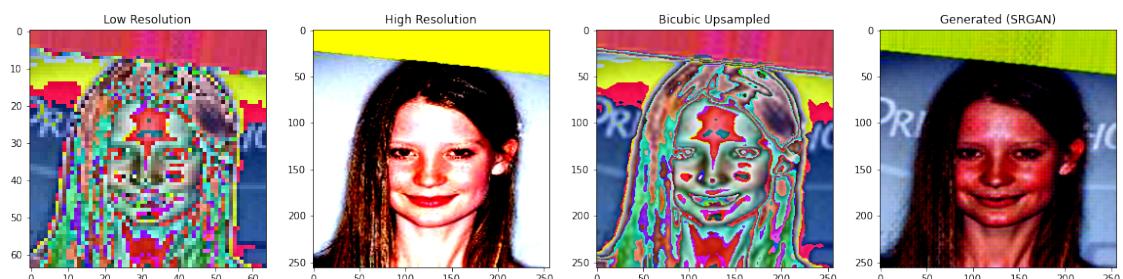
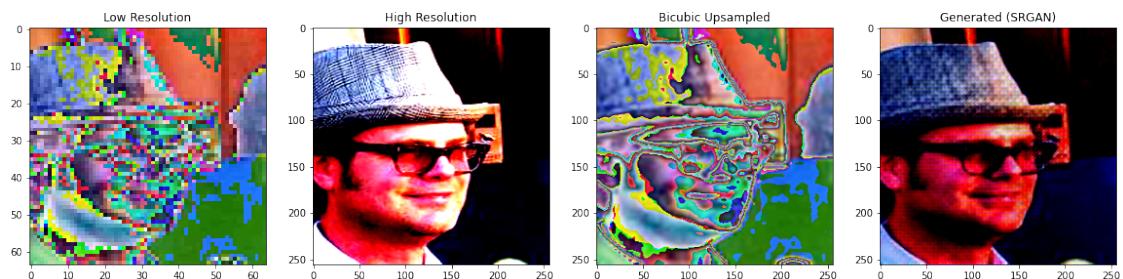
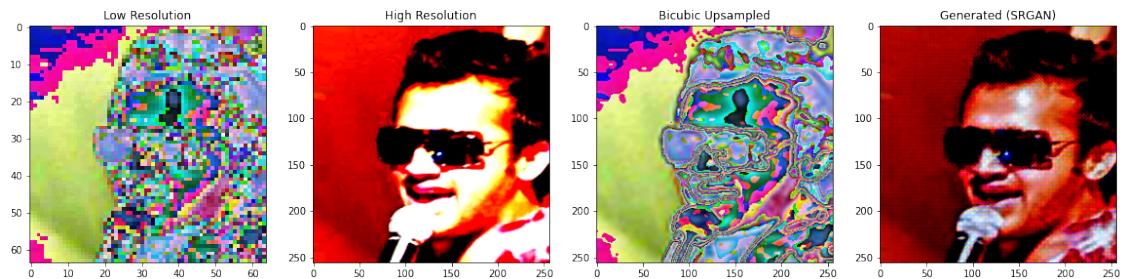
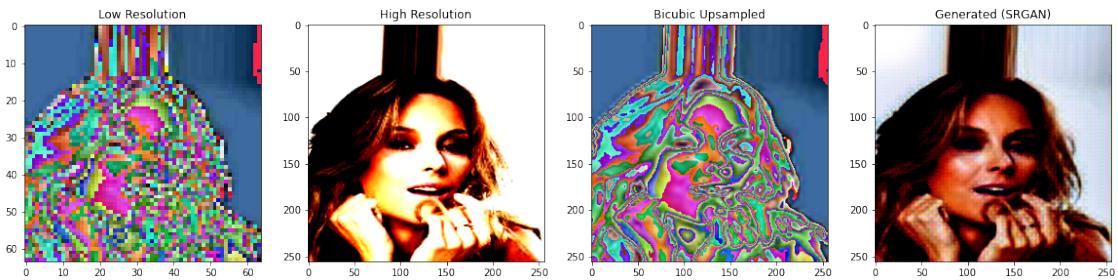


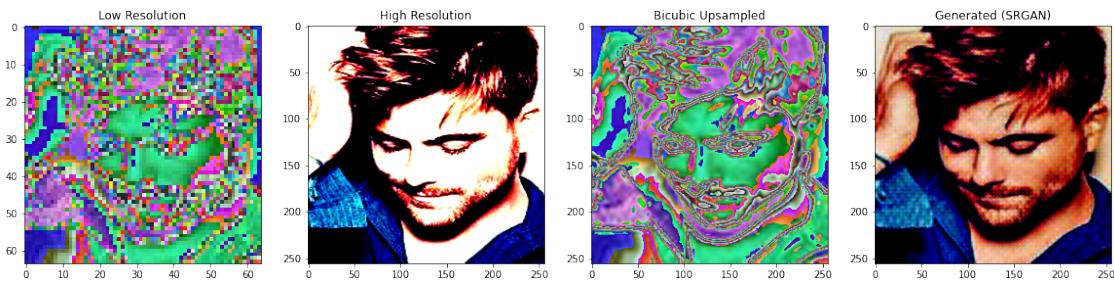
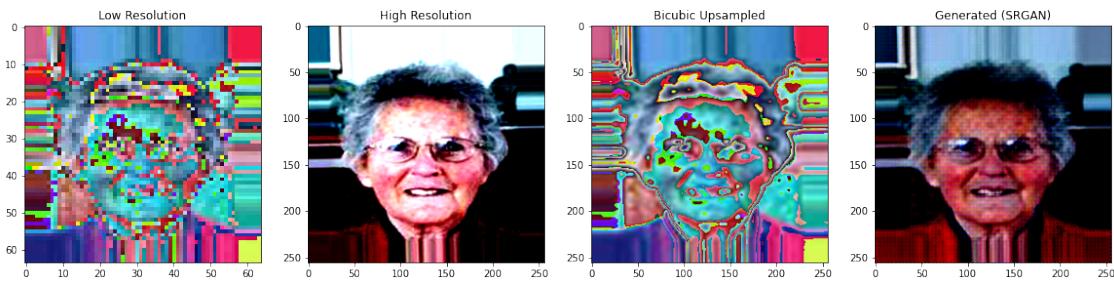
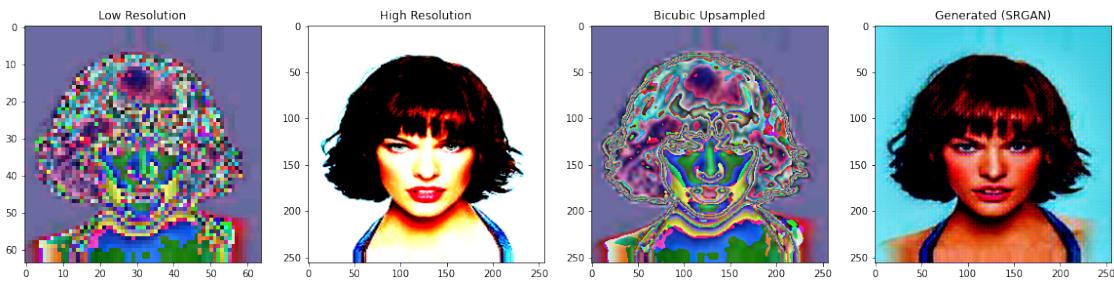
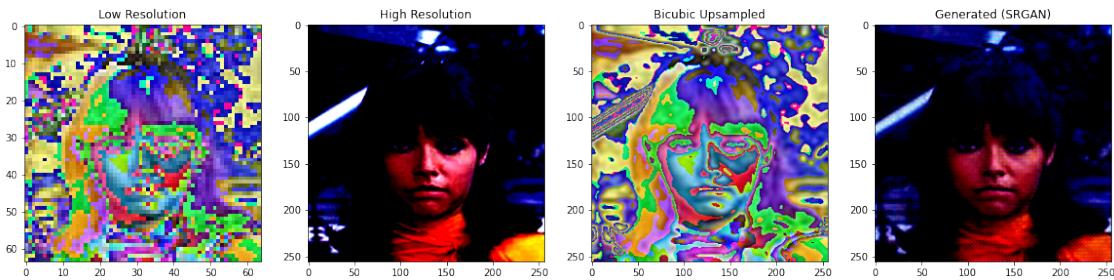


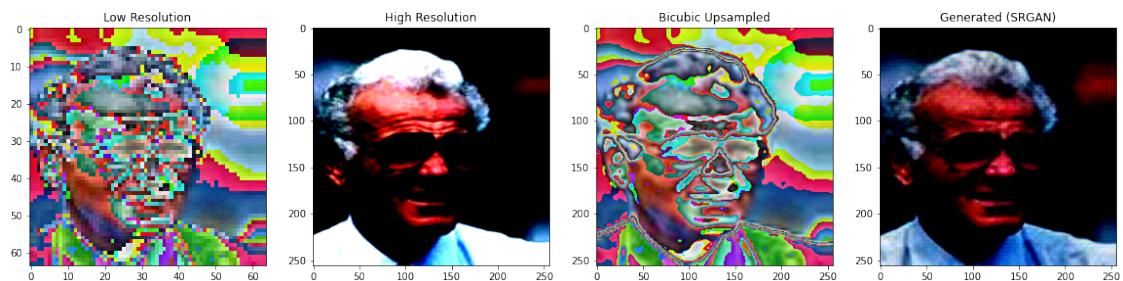
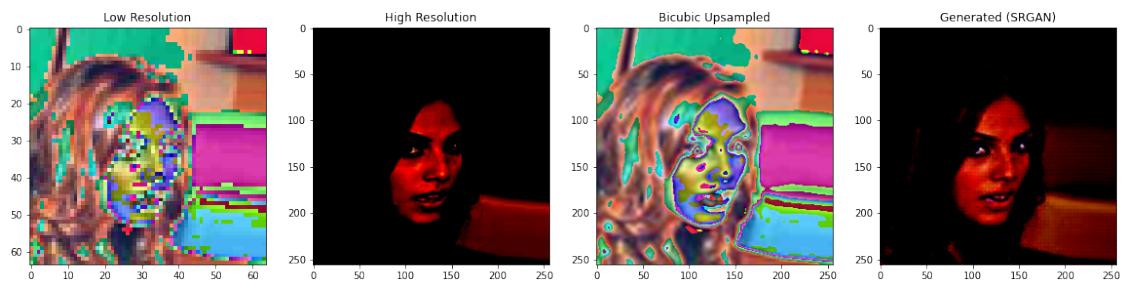
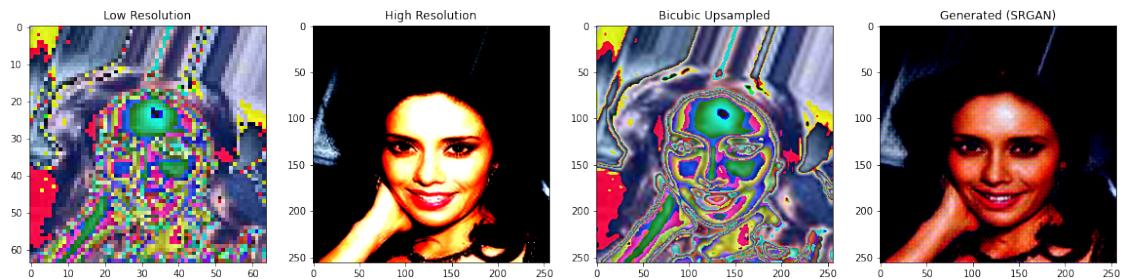
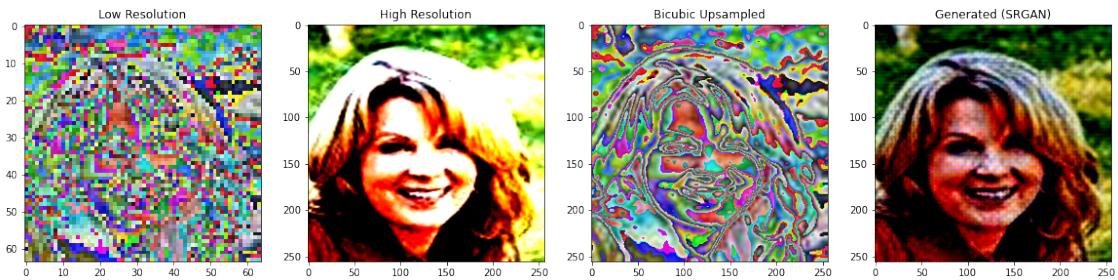


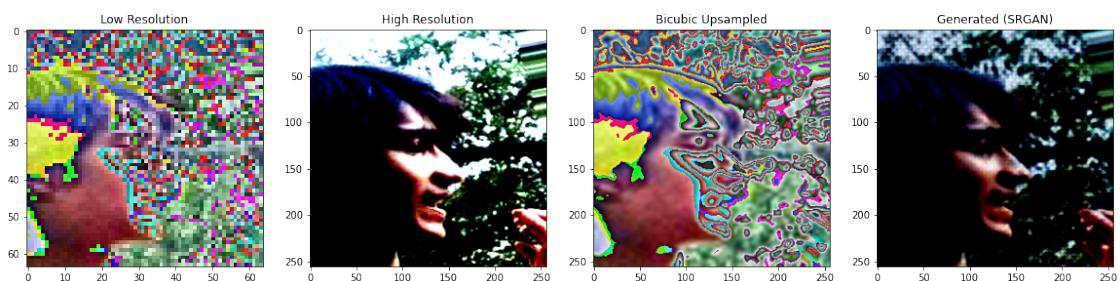
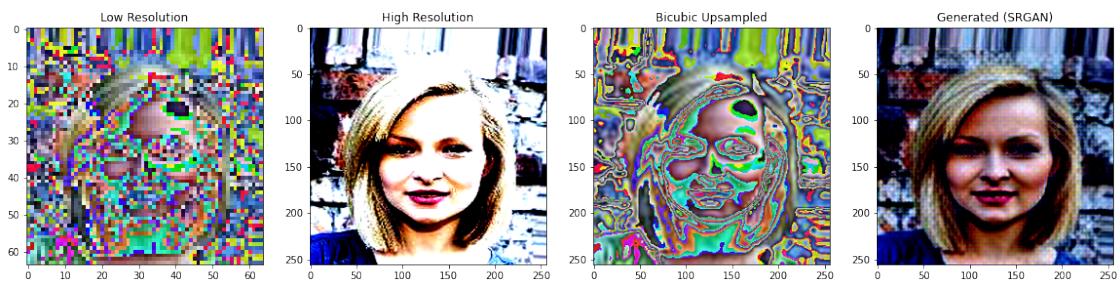
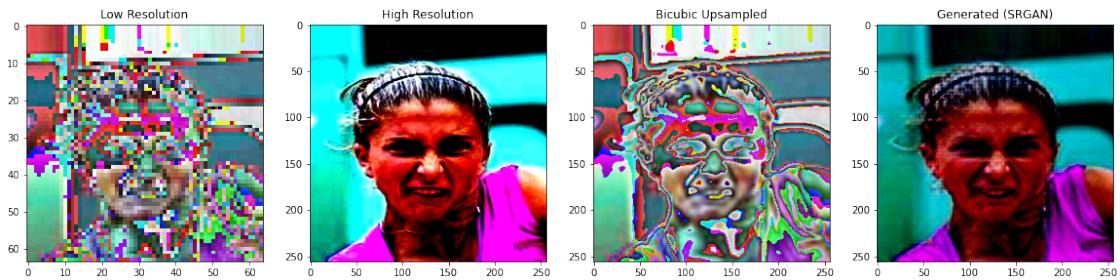




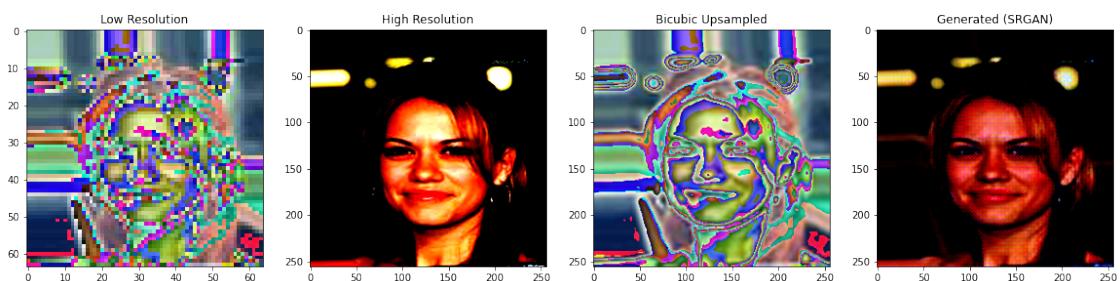


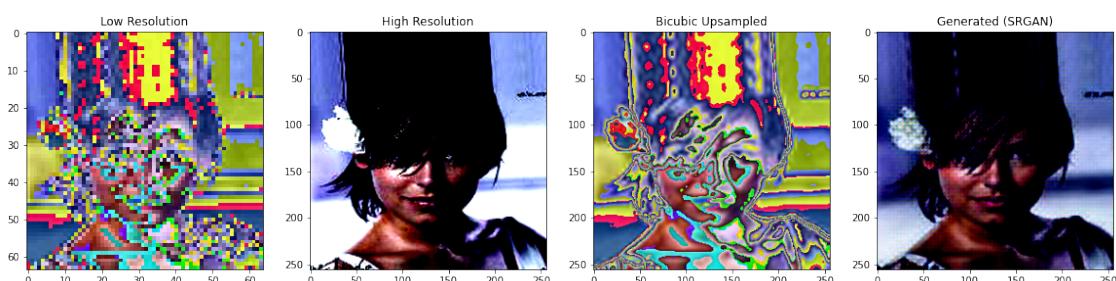
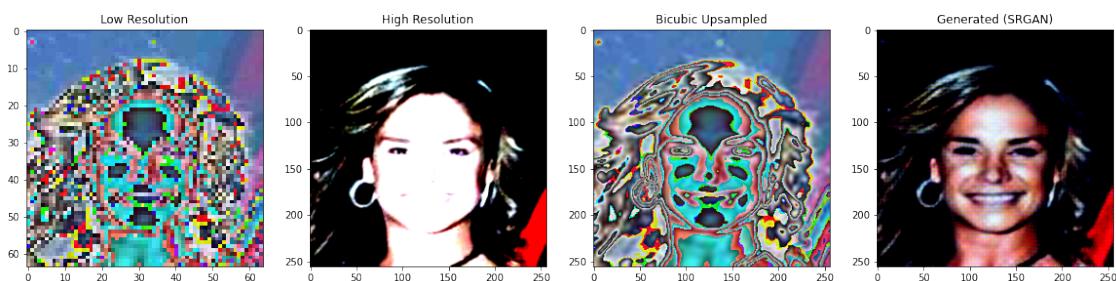
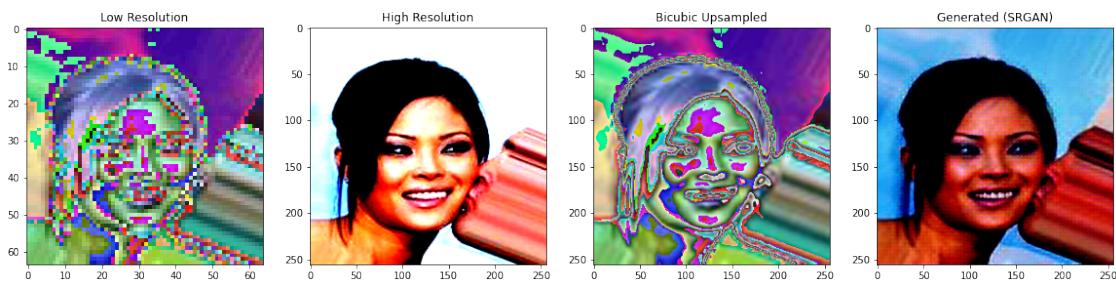
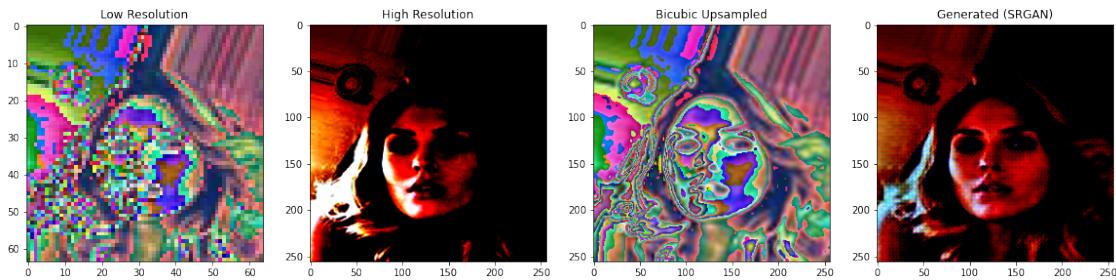


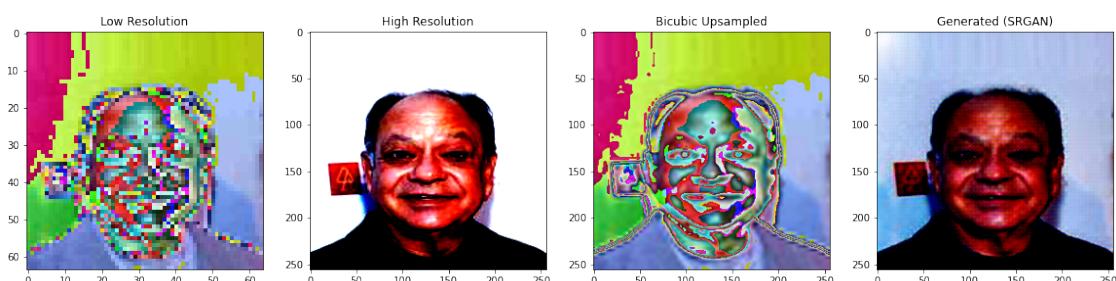
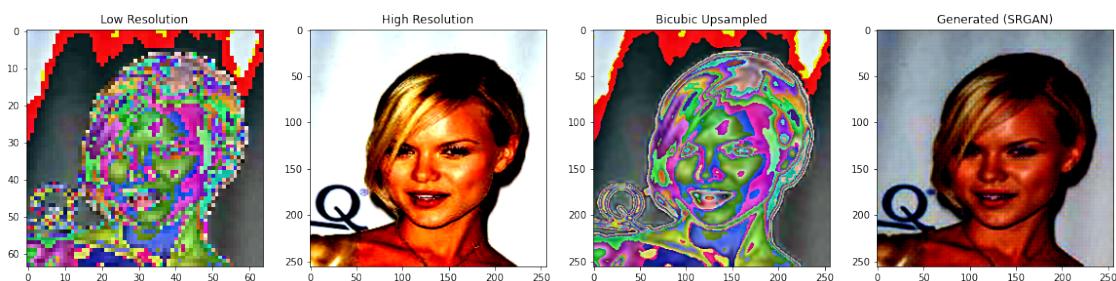
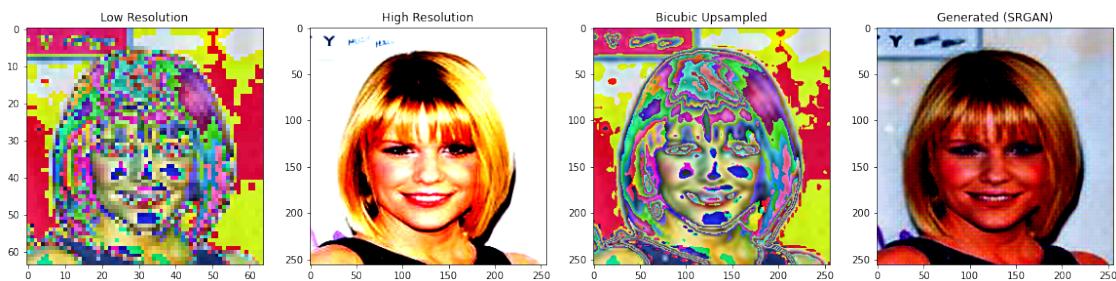
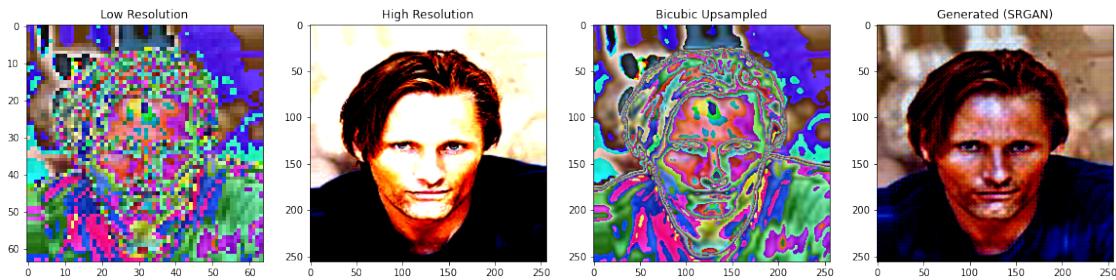


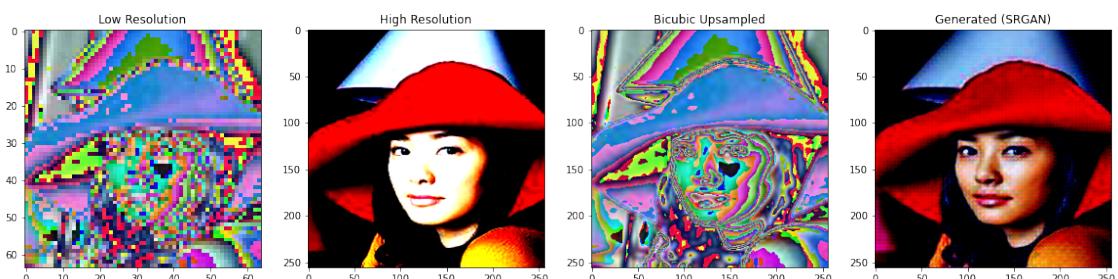
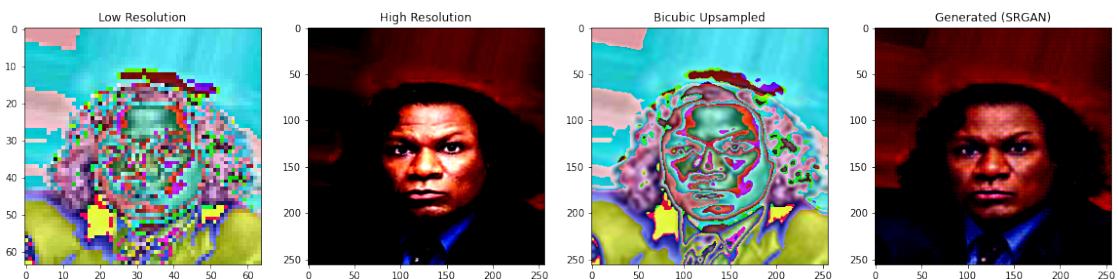
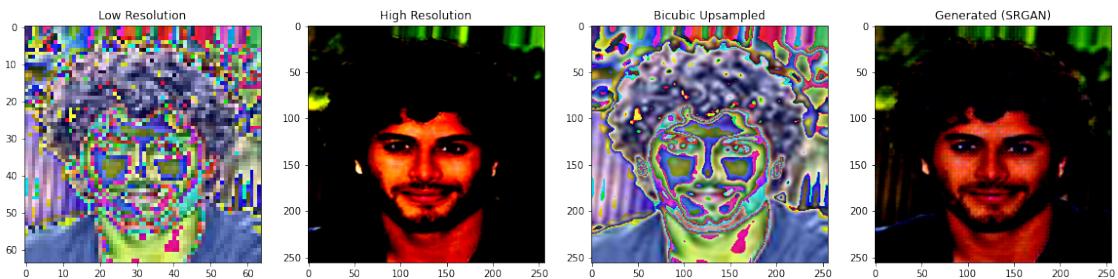
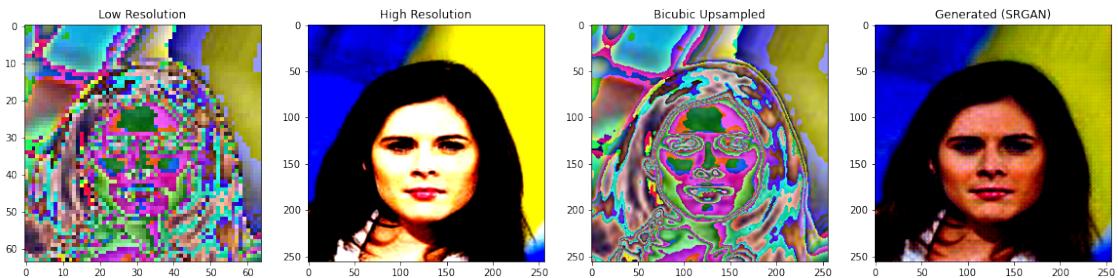


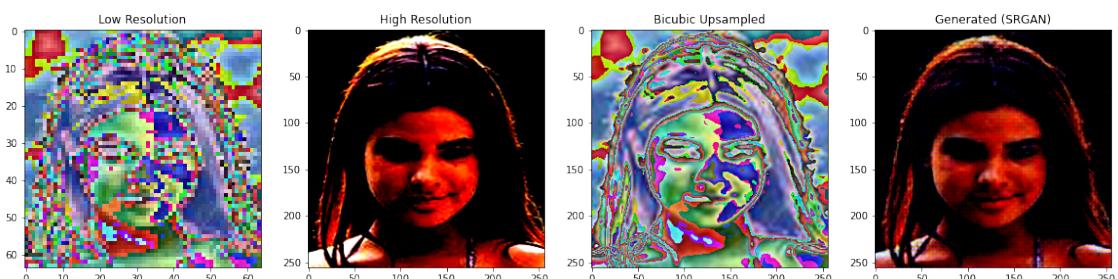
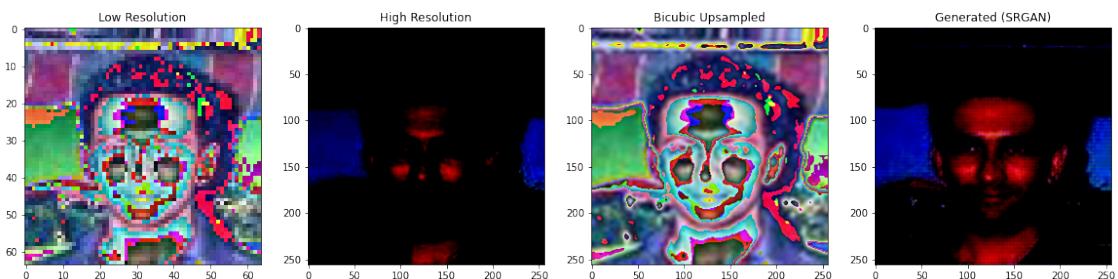
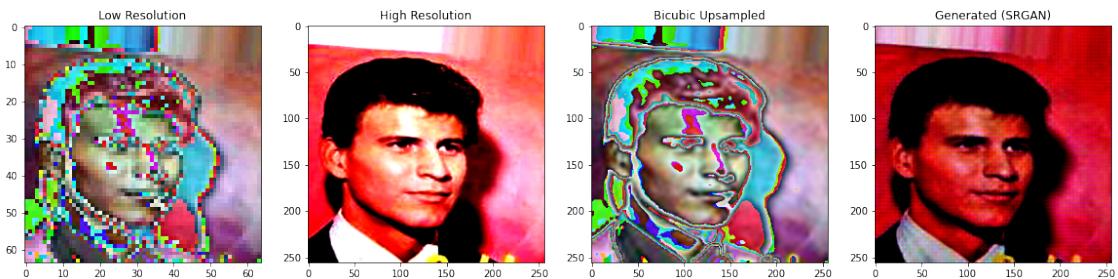
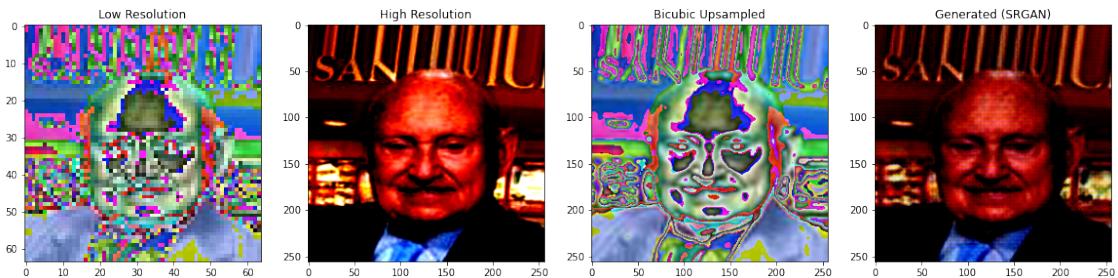
```
HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 1 ', max=338.  
~0, style=ProgressStyle(descrip...
```

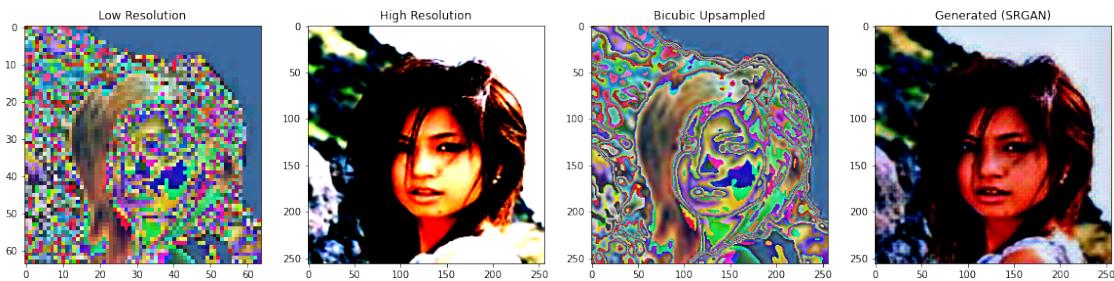
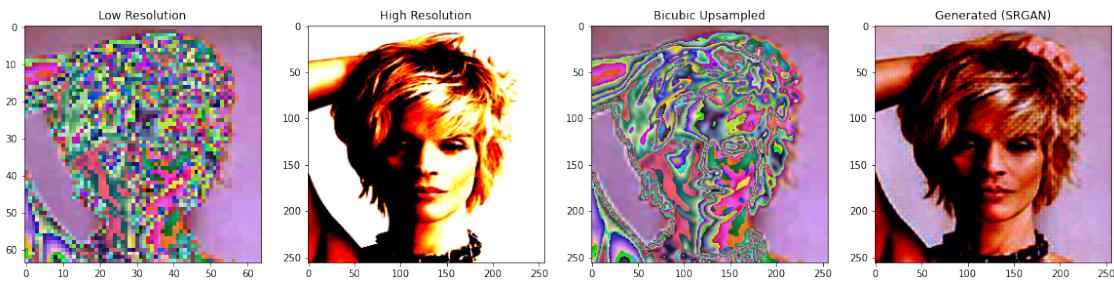
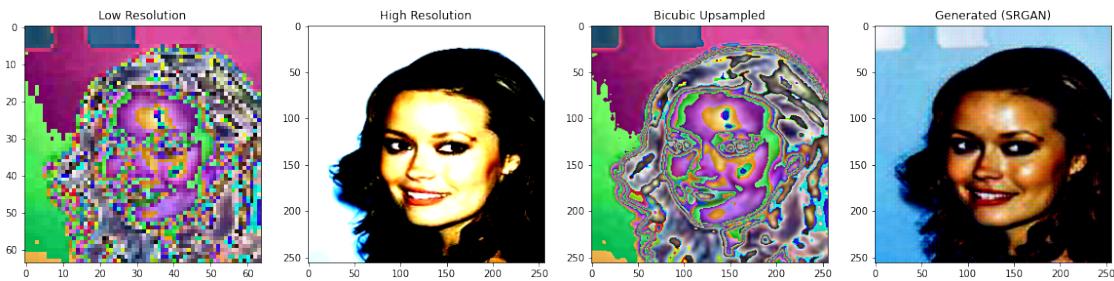
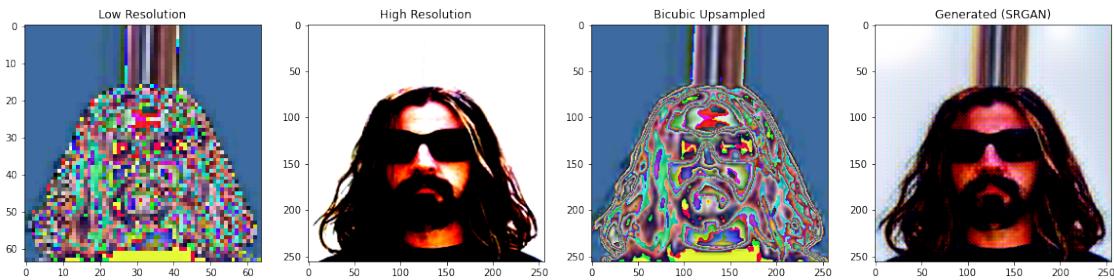


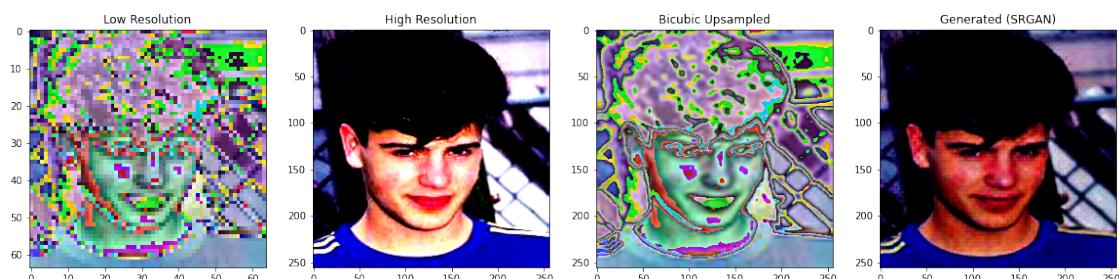
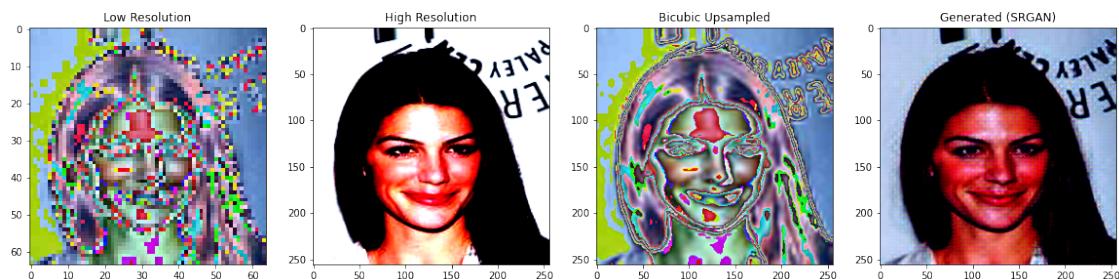
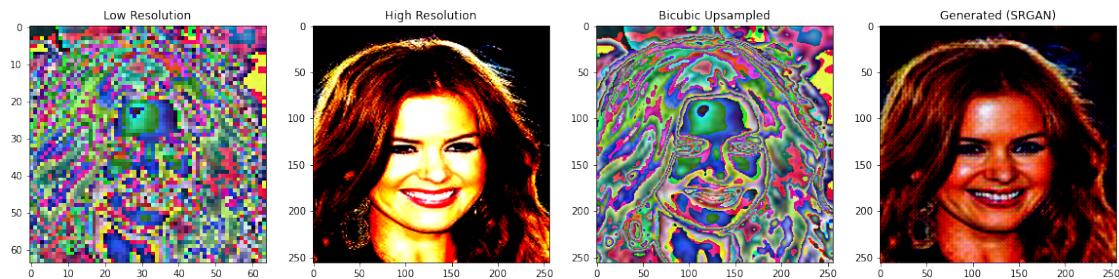
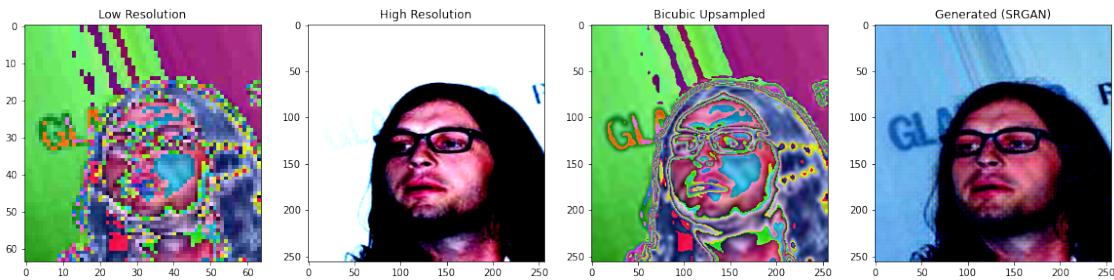


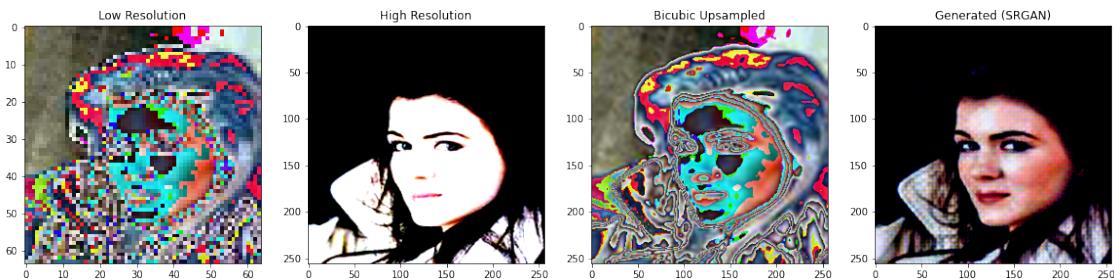
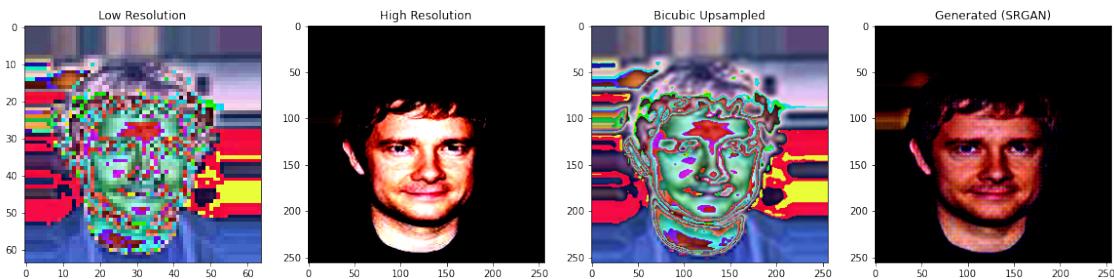
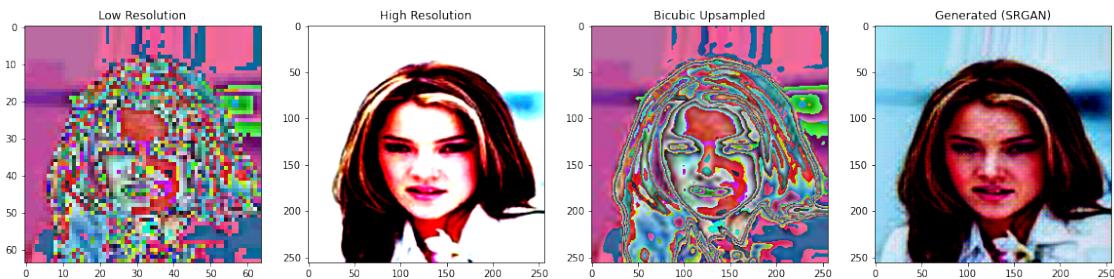
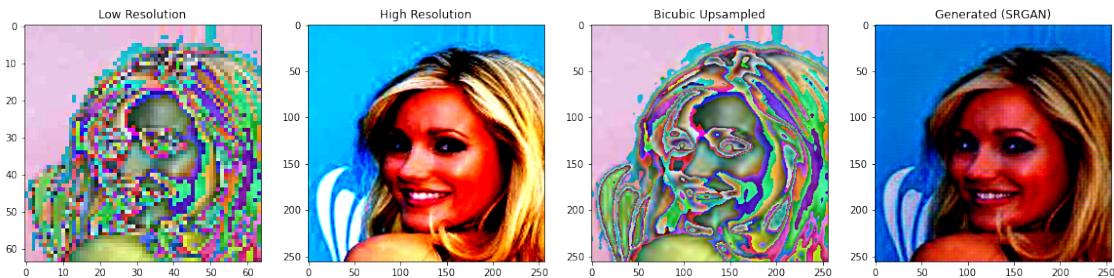












Epoch	Batch	PSNR	HR_Min	HR_Max	Gen_Min	Gen_Max
0	0	0	12.057614	0.0	1.0	0.0 1.000000
1	0	1	11.630188	0.0	1.0	0.0 1.000000

```

2      0      2  10.940614      0.0      1.0      0.0  0.999924
3      0      3  12.354712      0.0      1.0      0.0  1.000000
4      0      4  12.143303      0.0      1.0      0.0  1.000000
..
671     1    333  12.566619      0.0      1.0      0.0  1.000000
672     1    334  11.238737      0.0      1.0      0.0  1.000000
673     1    335  12.206124      0.0      1.0      0.0  0.999938
674     1    336  12.686282      0.0      1.0      0.0  1.000000
675     1    337  9.856307      0.0      1.0      0.0  1.000000

```

[676 rows x 7 columns]

```

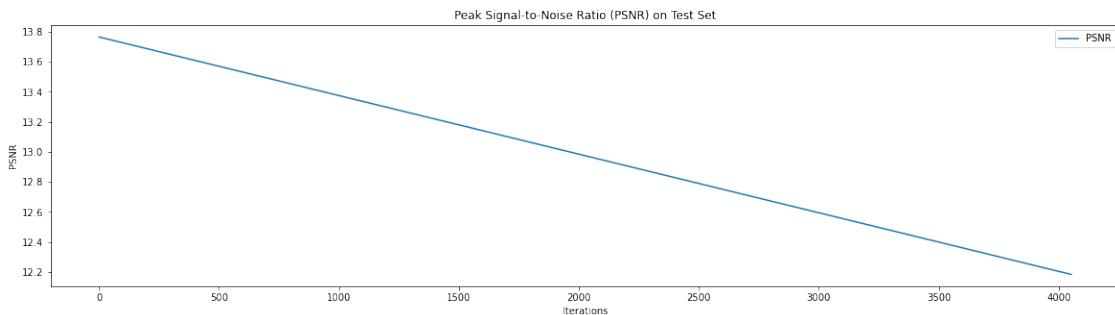
[54]: import matplotlib.pyplot as plt

# Ensure test_counter and psnr_values have the same length
min_length = min(len(test_counter), len(psnr_values))
test_counter = test_counter[:min_length]
psnr_values = psnr_values[:min_length]

# Adjust the figure size based on the total number of iterations
plt.figure(figsize=(20, 5))

plt.plot(test_counter, psnr_values, label='PSNR')
plt.legend()
plt.title('Peak Signal-to-Noise Ratio (PSNR) on Test Set')
plt.xlabel('Iterations')
plt.ylabel('PSNR')
plt.show()

```

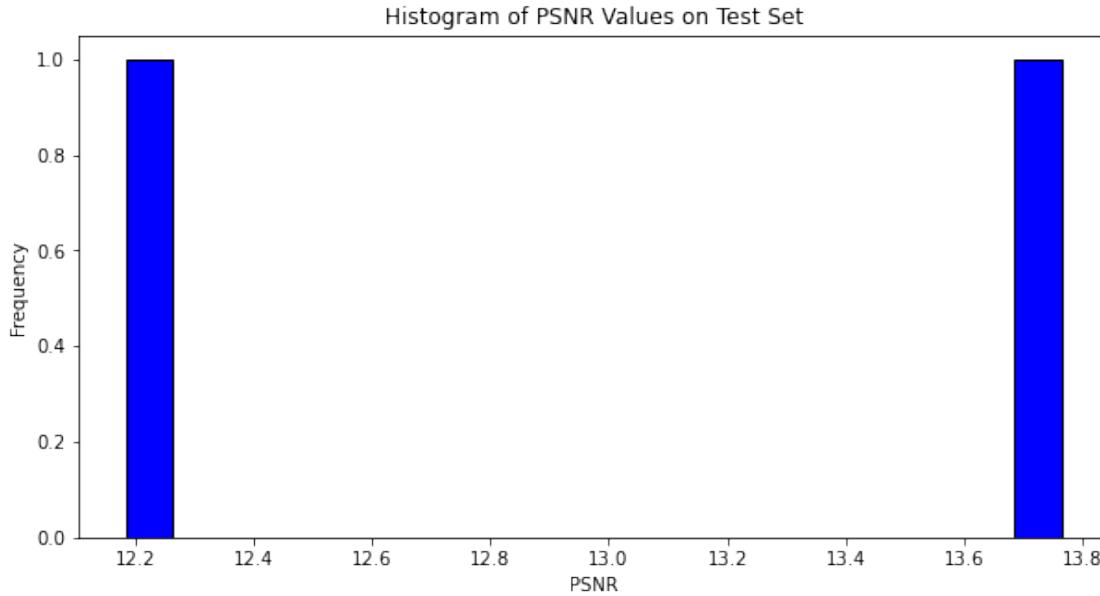


```

[55]: plt.figure(figsize=(10, 5))
plt.hist(psnr_values, bins=20, color='blue', edgecolor='black')
plt.title('Histogram of PSNR Values on Test Set')
plt.xlabel('PSNR')
plt.ylabel('Frequency')

```

```
plt.show()
```



```
[56]: from skimage.metrics import structural_similarity as ssim
import numpy as np

# ...

# Create lists to store SSIM values
ssim_values = []

for epoch in range(n_epochs):
    tqdm_bar = tqdm(test_dataloader, desc=f'Testing Epoch {epoch} ', ▾
        total=int(len(test_dataloader)))
    for batch_idx, imgs in enumerate(tqdm_bar):
        generator.eval(); discriminator.eval()
        # Configure model input
        imgs_lr = Variable(imgs["lr"].type(Tensor))
        imgs_hr = Variable(imgs["hr"].type(Tensor))
        # Adversarial ground truths
        valid = Variable(Tensor(np.ones((imgs_lr.size(0), *discriminator. ▾
            output_shape))), requires_grad=False)
        fake = Variable(Tensor(np.zeros((imgs_lr.size(0), *discriminator. ▾
            output_shape))), requires_grad=False)

        ### Eval Generator
        # Generate a high resolution image from low resolution input
        gen_hr = generator(imgs_lr)
```

```

# Ensure images are in the range [0, 1]
imgs_hr = imgs_hr.clamp(0, 1)
gen_hr = gen_hr.clamp(0, 1)

# Convert PyTorch tensors to NumPy arrays and change channel order if necessary
img1 = np.transpose(imgs_hr.cpu().detach().numpy(), (0, 2, 3, 1))[0]
img2 = np.transpose(gen_hr.cpu().detach().numpy(), (0, 2, 3, 1))[0]

# Calculate SSIM
ssim_value, _ = ssim(img1, img2, full=True, multichannel=True)

# Append SSIM value to the list
ssim_values.append(ssim_value)

tqdm_bar.set_postfix(gen_loss=loss_G.item(), disc_loss=loss_D.item(), ssim=ssim_value)

```

```
HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 0 ', max=338,
                             style=ProgressStyle(descrip...
```

```
HBox(children=(FloatProgress(value=0.0, description='Testing Epoch 1 ', max=338,
                             style=ProgressStyle(descrip...
```

[57]:

```

min_length_ssims = min(len(test_counter), len(ssim_values))
test_counter_ssims = test_counter[:min_length_ssims]
ssim_values = ssim_values[:min_length_ssims]

# Plotting SSIM values
plt.figure(figsize=(10, 5))
plt.plot(test_counter_ssims, ssim_values, label='SSIM')
plt.legend()
plt.title('Structural Similarity Index (SSIM) on Test Set')
plt.xlabel('Iterations')
plt.ylabel('SSIM')
plt.show()

```

